

Research and score anything, at scale, for free.

Your target accounts. Your docs. Your criteria.
All scored, sorted, and backed by proof.

- TRY WITH YOUR OWN LIST IN 30S

`free-fleet quickstart --demo` → no keys, no data leaves your machine

Automated Research Suite • MIT • runs anywhere

Accounts scored, sorted, and backed by proof.

The top 25 accounts ranked by fit score — with grounded evidence of their exact technical gaps.

ranked_target_accounts.csv			Scored & sorted by fit • Top 5 shown	
RANK	ACCOUNT	SCORE	IDENTIFIED GAP	GROUNDED EVIDENCE OF GAP
#1	stripe.com	98	Legacy billing migration to Kafka	" <i>lead migration off legacy v1 billing pipeline</i> "
#2	hyper_ai	94	50k QPS database latency wall	" <i>hitting latency limits at 50k QPS on self-hosted Postgres</i> "
#3	pinecone.io	91	Multi-tenant cluster search scaling	" <i>scaling vector search across multi-tenant clusters</i> "
#4	posthog.com	87	ClickHouse 2B+ event ingestion	" <i>optimizing ClickHouse to ingest 2B+ daily analytics events</i> "
#5	supabase.com	84	Distributed edge replication lag	" <i>building globally distributed real-time Edge Functions</i> "

PRIORITIZED PIPELINE

Scored & Sorted Accounts

Reps focus on top-scoring accounts first instead of guessing which cold accounts have active needs.

GROUNDED EVIDENCE

Documented Technical Gaps

Every qualification is backed by verbatim text proving the company’s real bottleneck.

The 4-Layer Filter.

From 1,000 down to the top 25.

Don't run one broad prompt. Compound layers of evidence, enriching only the survivors at each stage.

LAYER 1

1,000 → 600

Firmographic ICP Fit

Scans homepage & positioning. Immediately filters out wrong audience, B2C, and out-of-scope company types.

LAYER 2

600 → 150

Tech Stack & Architecture

Scans developer docs, API changelogs, and integrations. Validates tech compatibility with verbatim quotes.

LAYER 3

150 → 50

Hiring & Budget Signals

Scans active job descriptions and career pages. Pinpoints open roles citing target technologies or bottlenecks.

LAYER 4 • FINAL 25

50 → 25

Scoring & Gap Analysis

Ranks surviving accounts by ICP fit score (0–100) and extracts verbatim evidence of their core technical gaps.

Why this compounds so powerfully

By Layer 4, you don't just have account names. You have 25 top-scoring accounts with verbatim evidence proving their exact technical gaps and active bottlenecks.

Don't trust AI summaries.

Require exact quotes.

Every account qualification is grounded in verbatim evidence of their real technical gaps.

STEP 1 • PUBLIC JOB POST / CAREER PAGE

Role: Staff Infrastructure Engineer

“We are looking for a Staff Engineer to join our Platform team. In this role, you will be responsible for **leading the migration of our legacy billing service to Apache Kafka** and modernizing our distributed streaming architecture.”

STEP 2 • VERBATIM VERIFICATION GATE

✓ Substring Matched

IDENTIFIED GAP

Legacy Billing Migration to Kafka

ICP FIT SCORE

98 • Tier-1 Target Account

GROUND EVIDENCE

“leading the migration of our legacy billing service to Apache Kafka”

If the model hallucinates a gap or assumes an unverified bottleneck, the substring match fails and the account is discarded. Every qualified account is backed by documented proof.

PROOF OVER PREDICTIONS

Documented Gaps

Account qualification is anchored in an explicit bottleneck cited directly by the company.

ZERO HALLUCINATIONS

1-Click Verifiable Audit

Anyone on your team can verify the exact sentence on the live page proving why this account qualified.

Run it where you already chat.

Connect once, then score accounts and extract verified gaps in plain English.

1. CONNECT

Enable Tool

Add free-fleet to Claude, Codex, Cursor, Grokbot, or any MCP host.

2. DROP CSV

Upload Accounts

Drag & drop your list of companies or website URLs into chat.

3. CHAT

Ask in Plain English

Ask Claude to score accounts, sort them, and extract grounded evidence of technical gaps.

Claude / Codex / Cursor / Grokbot

YOU

“Triage **accounts.csv** to score the top accounts and extract grounded evidence of infrastructure bottlenecks and data pipeline migrations.”

CLAUDE

✓ Processed 100 accounts in 2 minutes • Scored and ranked top 25 accounts with grounded evidence of technical gaps • Exported **ranked_accounts.csv**

LOCAL & PRIVATE

Permissive MIT License

Open source and private. Account data and qualification criteria never leave your machine.

OPEN SOURCE

Star on GitHub ★

Quickstart guide and templates available at github.com/NatesVibeCode/free-fleet.