

2,500 → 25

in 25 min | \$0.00 API Spend

Open source research – every row includes the source quote

Starting set • example 2,500

2,500 input

Layer 1: Firmographic check 600

600 remaining

Layer 2: Tech stack notes 150

150 remaining

Layer 3: Hiring signals 50

50 remaining

LAYER 4 • FOUNDER NOTES

25 selected for review

→ 25 SELECTED

418 batches (15m) → 100 batches (4m)

Research suite filters 94% of raw noise before deep enrichment — via Claude/Cursor or CLI.

1 Layer 1: Firmographic Triage

418 Batches • 15 Mins

INPUT ACTION → EVIDENCE CHECK
SCRAPE_1 (Homepage & Pricing) → GATE_1 (Fit Score ≥ 70%)

FILTERED AT GATE 1
1,900 filtered
Not matching firmographic check

REMAINING
600 kept
Quote found on homepage/pricing

600 PASSED TO NEXT LAYER — EXAMPLE

2 Layer 2: Tech stack

100 Batches • 4 Mins

INPUT ACTION → EVIDENCE CHECK
SCRAPE_2 (Docs, APIs, Changelogs) → GATE_2 (Pain Score ≥ 75%)

FILTERED AT GATE 2
450 filtered
No relevant tech signal in docs

REMAINING
150 kept
Quote found in docs/changelog

25 batches (2m) → 12 batches (1m)

Research suite isolates hiring budget & founder bottleneck language — verbatim.

3

Layer 3: Hiring signals

25 Batches • 2 Mins

INPUT ACTION

SCRAPE_3 (Careers, Greenhouse, Lever)

→

EVIDENCE CHECK

GATE_3 (Open Reqs Matching Problem)

FILTERED AT GATE 3

100 filtered

No matching hiring signal

REMAINING

50 kept

Quote found in job post

4

Layer 4: Founder notes

12 Batches • 1 Min

CHECK

Founder posts / podcasts

→

KEEP IF

Quote describes bottleneck

FILTERED AT GATE 4

25 filtered

No relevant founder quote

REMAINING

25 kept

For review — with quotes

bash ~ live offline proof

\$ free-fleet quickstart --demo --run-id demo

✓ Deterministic demo pipeline • 0 API keys required • Validates quotes & schemas in <1s

Auto-reviewed against the source

If the model can't point to the exact quote, the row doesn't ship — we try the next lane.

\$ free-fleet export demo --format csv

3 of 25 verified records

ITEM_ID	PRIORITY	PRIMARY_QUOTE_TEXT
acme_cloud	HIGH	" scaling vector search across multi-tenant clusters is our #1 bottleneck this quarter"
stripe_billing	HIGH	"seeking Staff Engineer to lead migration off legacy v1 billing pipeline "
hyper_ai	HIGH	"currently using legacy self-hosted Postgres but hitting latency limits at 50k QPS "

~550

TOTAL BATCHES

\$0

TOTAL API COST

25 min

RUN TIME

25

TIER-1 ACCOUNTS

Run the Fleet on Your Machine

github.com/NatesVibeCode/free-fleet

STAR ON GITHUB ★

One command. No keys.

Anyone with Claude Desktop or Cursor can run the Research Suite — no CLI expertise needed.

bash — install to Claude/Cursor

```
$ pip install free-fleet
```

```
$ free-fleet mcp install --workspace-root "$PWD"
```

Auto-writes claude_desktop_config.json / mcp.json → restart app

```
$ free-fleet mcp install --dry-run --json # preview
```

```
$ free-fleet quickstart --demo --run-id demo # 0 keys, <1s proof
```



Or chat in Claude/Cursor — no terminal

In Claude: “Use free-fleet to triage `accounts.csv` for hiring intent — free-only”

→ Agent calls `free_fleet_run` → SQLite leases → verifies quotes → `free_fleet_export` → CSV